



Training. Adam, lr 10^{-4} , weight decay 10^{-6} , StepLR(4, $\gamma=0.6$), batch 2–4 (VRAM-bound; the reference used 8), early stopping on validation loss with patience 30. Best epoch lands at **4–6** of 35–36 run — these heads overfit early and the early stopping is doing real work. **Metric caveat.** The OKS column is a *mean* OKS per image with uniform $\kappa_k = 1/K$, not COCO keypoint AP; the two are not like-for-like and the apex column in particular must not be read as beating the reference.